

CSE598-Spring 2020

Assignment8.

Dense Matrix Multiplication Comparison: Hand-coded, Breeze, Spark

Project Report

Professor: Zou. Jia

Student Name: Weichi Zhao

Student ID: 1209692845

Email: wzhao42@asu.edu

- Hand-coded implementation can work: 3pt

```
235.72977312008226 252.88690305799346 257.7261071487286 256.654293408066 246.86907071034776 251.17513357449062 254.49413546911236 248.0
230.4289002129897 242.29005686522243 247.25507773282257 249.06769906153343 240.70191945895908 250.40399524715798 250.32381536288084 239
240.95843398908145 261.47284709334525 257.5733504044482 261.781634011792 247.19643842936205 257.1685109615922 258.00132479758514 248.78
243.62963995654408 252.58329046606318 256.68133073538394 261.74509140086184 249.41536232788698 254.19255374959178 250.69116160355716 23
242.76553440515463 248.75725109786702 253.27383707628775 251.64391934269221 249.10055646901674 253.44304956933368 255.47922933952424 24
245.07044026009856 259.10655210480837 204.69461295042015 272.41101841377721 256.4658037103749 270.0318456841251 260.247123597797094 257.06
239.8452382785641 258.0787028197392 254.21967419798935 260.30805857377166 247.07529167676158 251.67840879603443 253.47893846293607 245.
236.06452118511404 249.67978385540474 256.02131427127466 265.05490732318935 246.12903115406178 260.7819197078856 253.43159003827864 247
234.12371999704308 250.43064303316072 254.95244638594613 255.27314615227772 242.55734634252107 259.5859565204178 253.81302374009190 239
234.7777120729522 248.70030338675656 251.47811854192298 252.9043266344241 240.63775570088936 251.73461316409603 247.98216393615758 242.
232.26150951715087 239.14567620690883 248.41475537285557 258.29492461594305 241.7583591237724 252.16380374285438 249.66491171610448 234
241.55073170040326 259.66270267982793 264.58282812273615 263.1052790376752 252.5054750219284 263.8752025600239 258.10355490766375 248.2
239.33107369865417 250.14283216760518 254.96532008140574 255.51195467222541 246.8923654790311 260.78415847682066 253.0431003605605 243.
235.311338371044 254.53294955905838 255.85978531405745 261.54140091749018 251.47922766738043 261.6971834279119 255.96938345177654 242.99
247.44912944239914 260.35033953900005 256.7560826517732 262.8101538631769 254.1092721934339 261.9955091824311 260.0000541074917 244.434
245.01925140077324 251.846565844730813 259.33667405973456 262.16800250071177 250.71186319602978 260.47232928431305 252.3957798117815 243
233.19096892146595 247.14623339595994 247.778496047435 250.99765031892753 237.44290593197871 246.88260573755454 252.53004134264344 237.4
240.22077826042466 250.43502734775464 255.6260475218437 261.706105155756 241.4276478607329 252.70572459191314 252.40512122235206 242.94
239.9376293962724 247.32998503638285 248.10709471404859 253.49008409799892 243.04930055742403 253.61156207596512 252.70273261740603 244
228.650260391992 241.67617303333487 242.86297587217302 245.12707011678023 232.3057151401967 244.26912714016692 241.14450992440655 236.5
243.15442135511458 254.20445442560424 253.56597307611696 259.233900608311096 244.8749298400432 258.7292613599439 253.24974785798 238.885
243.09740968108237 256.53095296741726 258.55590910333197 264.1553361467152 251.54990487039515 260.17334121455355 254.2256324476929 247.
```

```
Process time of Dense Matrix Multiplication using Hand-coded scala: 13.006358609 second(s)
Average process time of Dense Matrix Multiplication using Hand-coded scala: 0.13006358609 second(s)
```

○ Process finished with exit code 0

- Spark implementation can work: 3pt

```
242.9867376894244 240.5825547266446 251.5181927924703 ...
244.59990744725627 248.676496028416 252.426770050094335 ...
249.791275652174 252.19696340090385 252.6966172564088 ...
251.37674250239166 250.5075413546023 253.8527149068731 ...
240.72794601938102 238.75058788896465 245.22262038233654 ...
244.80146781625157 247.97261319631917 245.25700409123774 ...
253.66813054003654 254.76470681891894 250.64514409864813 ...
254.53072287553647 253.97937413250312 252.21323857561384 ...
246.22275005598607 248.51237083348016 247.33250223462954 ...
249.21138910798072 255.16426518557256 253.57005803977256 ...
254.77628455634294 252.50901513958195 255.2510632554872 ...
248.51599246655755 250.70379556191213 255.3546740202075 ...
256.51240700257995 257.94929039665343 260.48475733108137 ...
249.9147221042649 254.32735329094365 252.16173821138986 ...
246.6484393111493 251.42728582510864 250.60222783424467 ...
253.7471871547881 254.6162582264906 254.8861678984912 ...
244.34703051496527 250.08128592983925 247.62279709881284 ...
255.244789779947 256.98281357992556 245.9889723298474 ...
255.20057054615018 258.78903137323107 256.0026070271346 ...
249.59119300484292 250.57297289360707 255.84107533436747 ...
248.51382372999524 257.699775789721 251.26472531070414 ...
248.59191925097082 251.00159488200993 246.40311828530383 ...
... (1000 total))
```

```
ProcessTime of Dense Matrix Multiplication using Spark: 1.058135671 second(s)
Average process time of Dense Matrix Multiplication using Spark: 0.01058135671 second(s)
```

```
20/04/16 23:50:12 INFO SparkUI: Stopped Spark web UI at http://192.168.223.128:4040
20/04/16 23:50:12 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
```

○

- Breeze implementation can work: 3pt

```
251.47878185462542 248.45668842400116 242.4007610987665 ...
255.95707092243714 251.88585302239235 252.51037737355273 ...
255.27698838151917 247.41248406976504 243.37714312695772 ...
253.43775358007902 245.1227518748211 250.0506128129136 ...
255.1606269140931 253.04055328316802 246.34507335091405 ...
268.52641341175837 262.49299811366234 253.43006237993893 ...
267.8559403924122 259.7478200831416 253.6050107267238 ...
259.9930826310499 256.7047605782091 250.5999059102858 ...
262.8609588144696 258.3745854931106 252.73965000119023 ...
252.15131650405203 250.04963522881226 243.42976581202626 ...
258.1494096246601 254.46997393968593 251.28330647475195 ...
253.94264678198706 249.39255465961588 239.80033725055338 ...
267.40516660083614 258.56252507479746 252.31578805770404 ...
251.59081614303557 248.7537676435308 243.6728022004324 ...
244.76233122606087 239.5275810517491 234.03903616890918 ...
253.1397209742074 250.16210529512367 248.69600282995245 ...
256.4332589023562 251.8272970018962 250.75567320037536 ...
258.0267958609733 254.02456443772815 245.8400532306435 ...
249.58064950118777 243.75078805320006 233.651021320769 ...
251.13163426374805 247.2972173898635 243.80001140425396 ...
253.40796205432406 251.1819784588757 246.34594073197607 ...
263.6845722642794 258.4053600827644 247.67560186146596 ...
... (1000 total)
```

```
ProcessTime of Dense Matrix Multiplication using breeze: 1.02007906 second(s)
Average process time of Dense Matrix Multiplication using breeze: 0.0102007906 second(s)
```

○ Process finished with exit code 0

○

- **In the report, the machine hardware configuration is listed: 3pt**
 - **Framework:** VMware
 - **Virtualized System:** Ubuntu 18.04 LTS
 - **Virtualized RAM:** 4096 MB
 - **Virtualized CPU cores:** 4
- **In the report, measured time for the three implementations are listed and compared in a way that is easy to understand: 4pt**
 - Hand-coded Scala code for executing 1000x1000 dense matrix multiplication for N times and get the average time for executing one matrix multiplication.
 - **The number of times:**
 - 1000 times
 - **Processing time:**
 - 13.606358689 seconds
 - **Average time for executing one matrix multiplication:**
 - 0.13606358689 seconds
 - Spark code for executing 1000x1000 local dense matrix multiplication for N times and get the average time for executing one matrix multiplication.
 - **The number of times:**
 - 1000 times
 - **Processing Time:**
 - 1.058135671 seconds
 - **Average time for executing one matrix multiplication:**
 - 0.01058135671 seconds
 - Scala code using Breeze for executing 1000x1000 local dense matrix multiplication for N times and get the average time for executing one matrix multiplication.
 - **The number of times:**
 - 1000 times
 - **Processing Time:**
 - 1.02007906 seconds
 - **Average time for executing one matrix multiplication:**
 - 0.0102007906 seconds
 - **Comparison(processing time):**
 - Hand-coded Scala code > Spark code \approx Scala code using Breeze
 - Hand-coded Scala code took the longest time to finish matrix multiplication, Spark code and Scala code using Breeze have a similar processing time.

- **In the report, observations and things learned from the implementation and performance comparison are summarized: 4pt**
 - **Observation:** After implementing the three different approaches for matrix multiplication, I noticed that different procedures of mathematical calculation could result in significant performance differences. I also notice that Scala is a powerful OO programming language and Spark provides a human-kind interface for submitting the tasks.
 - **Things Learned:** During this assignment, I practiced the setup of the SBT environment, basic Scala programming ability, and Spark hands-on implementation for submitting the task to the local cluster. I learned how to hand-code the simple matrix multiplication by using Scala, and how to utilize the library like breeze and Spark-Mllib to create the matrix and implement matrix multiplication. And I also learned how to measure and evaluate the performance of the program in Scala and Spark.
 - **Performance Comparison:** While the implementation of matrix multiplication in three different approaches above, the results of processing time under the same hardware environment are different. The implementation of hand-coded in Scala took the longest time which is around 14 seconds and it took 0.13606358689 seconds to finish one matrix multiplication. Both performances of Scala code using breeze and Spark code are similar which is around 1 second and it took 0.01 second to do one matrix multiplication, but the breeze implementation is slightly faster than Spark implementation.